

Bracketing hyperoperations: a rank dichotomy for the Tamari order

Toshihisa Urahigashi

Independent researcher, Nara, Japan

ORCID: 0009-0004-0460-6242

Abstract

Let H_r denote the Goodstein hyperoperation hierarchy, so that H_1 is addition, H_2 multiplication, H_3 exponentiation and H_4 tetration. Since H_r is non-associative for $r \geq 3$, the value of a bracketed expression $H_r(a, \dots, a)$ depends on the bracketing, i.e. on a binary tree. We study how the value varies over the Tamari lattice of such trees.

We prove that a *semi-associative law* runs through the whole hierarchy: for every rank $r \geq 3$, all $x, y \geq 2$ and all $z \geq 0$,

$$H_r(H_r(x, y), z) \leq H_r(x, H_r(y, z)),$$

so a right rotation at any node never decreases the value and evaluation is order-preserving from the Tamari lattice to $(\mathbb{N}, \leq)$; the left comb is a minimiser and the right comb, whose value is $H_{r+1}(a, n)$, a maximiser.

Our main result is that the *degeneracy* of this law disappears at rank 4. We classify equality completely. For $x, y \geq 2$ and $z \geq 0$:

$$r = 3 : \quad \text{equality} \iff z = 1 \text{ or } (y, z) = (2, 2); \quad r \geq 4 : \quad \text{equality} \iff z = 1.$$

Thus rank 3 carries exactly one non-trivial *equality pattern* — already visible for three leaves, where $a = 2$ makes the single cover $(aa)a < a(aa)$ an equality, both sides being 16, and producing several collisions by four leaves — and from rank 4 on the evaluation is strictly increasing along *every* such cover (Theorem 1.4). (At rank 3 with $a = 2$ and $n = 4$ several trees share a value; what is unique is the non-trivial equality pattern $(y, z) = (2, 2)$ of the local law itself.) The rank-3 equality is produced by the collapsing law $(x^y)^z = x^{yz}$, which reduces the rank-3 inequality to $yz \leq y^z$; the classification shows that no equality of this kind survives at rank 4 or above.

The result is not an asymptotic triviality. The two sides compare a growing *first* argument against a growing *second* one, and small arguments do not blow up with the rank: $H_r(2, 2) = 4$ for every $r \geq 2$. Moreover the hypothesis $y \geq 2$ is sharp as a uniform hypothesis: for every $r \geq 3$ and all $x, z \geq 2$, taking $y = 1$ reverses the inequality.

The monotonicity and strictness statements are proved for arbitrary leaf values ≥ 2 , not only for a single repeated one. We stress what is *not* new: at rank 3 the monotonicity direction is already known, in a stronger form, from a 2018 MathOverflow discussion of exponentiation, and we do not claim it. What we add is the extension of that direction to every rank of the hierarchy, the classification of the equality cases, the strictness from rank 4 on, and the formalization.

Every numbered statement of the paper that asserts an inequality or an equality about H_r or about the evaluation of trees is formalized and machine-checked in Lean 4; the classical structure of the Tamari lattice itself (that it is a lattice with the two combs as its extreme elements) is used as background and is not formalized. The development contains no `sorry`, every theorem is audited with `#print axioms`, and no declaration depends on `Classical.choice`.

1 Introduction

Exponentiation is not associative, and it is classical that among all bracketings of $a^{a^{\dots^a}}$ with n occurrences of $a \geq 2$ the right-associated one is the largest; its value is the tetration $a \uparrow\uparrow n$. Orders on such bracketings have been discussed before — see Section 9 — but we are not aware of a treatment of the *Tamari* order on them, nor of the behaviour of that order further up the hyperoperation hierarchy.

Two strands of existing work bracket the question. On the one hand, the *associative spectrum* of a binary operation, introduced by Csákány and Waldhauser [1], measures non-associativity by counting how many distinct values the C_{n-1} bracketings can take; they prove in particular that exponentiation is *Catalan*, i.e. distinct bracketings induce distinct operations. That line of work counts values. On the other hand, the set of bracketings carries the Tamari lattice [12], generated by right rotations, whose proof-theoretic content was isolated by Zeilberger [13] as a semi-associative law. That line of work orders bracketings. The present paper connects the two: we show that for hyperoperations the Tamari order is *respected* by evaluation, and we locate precisely where the correspondence degenerates.

Results

Throughout, H_r is the hyperoperation hierarchy fixed in Section 2, with $H_1(a, b) = a + b$, $H_2(a, b) = ab$, $H_3(a, b) = a^b$ and $H_4(a, b) = a \uparrow\uparrow b$.

Theorem 1.1 (Rank 3, complete solution). *Let $x \geq 1$ and $y \geq 2$. Then $(x^y)^z \leq x^{(y^z)}$ for every $z \geq 0$. If moreover $x \geq 2$, equality holds if and only if $z = 1$ or $(y, z) = (2, 2)$. The hypothesis $y \geq 2$ cannot be dropped: for $y = 1$ and all $x, z \geq 2$ the inequality fails.*

Theorem 1.2 (Rank 4). *Let $x \geq 2$ and $y \geq 2$. Then $(x \uparrow\uparrow y) \uparrow\uparrow z \leq x \uparrow\uparrow (y \uparrow\uparrow z)$ for every $z \geq 0$. If moreover $z \geq 2$, the inequality is strict.*

Proposition 1.3 (Tamari monotonicity). *Let $a \geq 2$ and $r \geq 3$. If $T \triangleleft U$ is a Tamari cover, i.e. U arises from T by a right rotation at some node, then $\text{ev}_{r,a}(T) \leq \text{ev}_{r,a}(U)$. Consequently $\text{ev}_{r,a}$ is order-preserving for the Tamari order. Moreover, for the right comb R_n with n leaves, $\text{ev}_{r,a}(R_n) = H_{r+1}(a, n)$.*

Theorem 1.4 (Strict Tamari monotonicity from rank four). *Let $a \geq 2$ and $r \geq 4$. If $T \triangleleft U$ is a Tamari cover, then*

$$\text{ev}_{r,a}(T) < \text{ev}_{r,a}(U).$$

Thus for $r \geq 4$ the evaluation is strictly increasing along every cover of the Tamari lattice, not merely along root rotations.

Corollary 1.5 (Rank dichotomy). *For $a \geq 2$ the evaluation $\text{ev}_{r,a}$ is order-preserving for the Tamari order at every rank $r \geq 3$ (Proposition 1.3), and*

- *at rank 3 strictness can fail, but only for $a = 2$: the cover $((aa)a)a \triangleleft (aa)(aa)$ carries equal values, both 256, when $a = 2$;*
- *at every rank $r \geq 4$ it is strict on covers (Theorem 1.4).*

In fact $(r, a) = (3, 2)$ is the only exception: for every other pair with $r \geq 3$ and $a \geq 2$ every cover is strict (Corollary 7.3).

Theorem 1.6 (All ranks). *For every $r \geq 3$ and all $x \geq 2$, $y \geq 2$ and $z \geq 0$,*

$$H_r(H_r(x, y), z) \leq H_r(x, H_r(y, z)).$$

Theorem 1.7 (Complete equality classification). *Let $x \geq 2$ and $y \geq 2$.*

1. *For every $r \geq 2$ the two sides are equal when $z = 1$; in particular equality holds in Theorem 1.6 for every $r \geq 3$.*
2. *For $r = 3$ the only further equality is $(y, z) = (2, 2)$.*
3. *For $r \geq 4$ and $z \geq 2$ the inequality is strict.*
4. *For every $r \geq 3$ the inequality is strict when $z = 0$, since $H_r(H_r(x, y), 0) = 1 < x = H_r(x, H_r(y, 0))$.*

Items (1), (3) and (4) together give: for $r \geq 4$ equality holds if and only if $z = 1$.

Theorems 1.6 and 1.7 are proved in Section 6 by inductions on the rank whose engine is the identity-like inequality (SUM_r) below; the case $r = 3$ is Theorem 1.1 and the case $r = 4$ can also be obtained directly, as in Section 4.

Remark 1.8 (The statement is not an asymptotic triviality). Two features are worth isolating. First, (HT_r) does not follow from monotonicity alone: on the left the *first* argument grows, $x \mapsto H_r(x, y) > x$, while on the right the *second* argument grows, $z \mapsto H_r(y, z)$, so the two coordinates compete. Second, high rank does not make everything large: since $H_r(a, 1) = a$ for $r \geq 2$, one gets

$$H_r(2, 2) = H_{r-1}(2, H_r(2, 1)) = H_{r-1}(2, 2) = \cdots = H_2(2, 2) = 4$$

for every $r \geq 2$. A degenerate corner therefore survives at every rank, and it is exactly this corner that carries the rank-3 equality of Theorem 1.7(2) and that has to be treated separately in the proof of Theorem 1.7(3). Neither observation is by itself a proof that the statement is hard; what makes it worth isolating is the combination of the two with the fact that a single induction carries the law through every rank, that the degeneracy disappears exactly at rank 4, and that the equality set can be described completely.

Formalization

Every numbered statement of this paper that asserts an inequality or an equality about H_r , or about the evaluation of trees, is proved in Lean 4 on top of Mathlib; the classical structure of the Tamari lattice is used as background and is not formalized here. Section 8 makes the division precise, statement by statement. The development defines the hyperoperation by the usual double recursion and *verifies the convention* by proving $H_1 = +$, $H_2 = \cdot$, $H_3 = ()^{()}$; this matters because the literature on hyperoperations uses several incompatible indexings and initial values. Section 8 describes the development.

2 Preliminaries

The literature on hyperoperations uses several mutually incompatible indexings and initial values. Throughout, H_r means the Goodstein-style convention fixed in Definition 2.1; where a statement depends on the value at $b = 0$ — as the equality classification does at $z = 0$ — this is noted.

Definition 2.1 (Hyperoperation [7]). Define $H : \mathbb{N} \times \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}$ by

$$\begin{aligned} H_0(a, b) &= b + 1, & H_1(a, 0) &= a, & H_2(a, 0) &= 0, & H_{r+3}(a, 0) &= 1, \\ H_{r+1}(a, b + 1) &= H_r(a, H_{r+1}(a, b)). \end{aligned}$$

Lemma 2.2. *For all a, b we have $H_1(a, b) = a + b$, $H_2(a, b) = ab$ and $H_3(a, b) = a^b$. Moreover $H_r(a, 1) = a$ for $r \geq 2$.*

We write $a \uparrow\uparrow b := H_4(a, b)$, so $a \uparrow\uparrow 0 = 1$ and $a \uparrow\uparrow (b + 1) = a^{(a \uparrow\uparrow b)}$.

Definition 2.3 (Bracketings and evaluation). Let $\mathcal{T}_n$ denote the set of binary trees with n leaves. For $r, a \in \mathbb{N}$ define $\text{ev}_{r,a} : \mathcal{T}_n \rightarrow \mathbb{N}$ by $\text{ev}_{r,a}(\text{leaf}) = a$ and $\text{ev}_{r,a}(T \cdot U) = H_r(\text{ev}_{r,a}(T), \text{ev}_{r,a}(U))$.

Definition 2.4 (Right rotation). The *right rotation at the root* sends $((A \cdot B) \cdot C)$ to $(A \cdot (B \cdot C))$. The Tamari order on $\mathcal{T}_n$ is generated by right rotations performed at arbitrary nodes; that this order is a lattice, with the left comb as least and the right comb as greatest element, is the theorem of Friedman and Tamari [2], whose title already names the relation a *loi demi-associative* [12, 13].

Thus the inequality

$$H_r(H_r(x, y), z) \leq H_r(x, H_r(y, z)) \quad (\text{HT}_r)$$

is exactly the assertion that a right rotation, applied with x, y, z the values of the three subtrees, does not decrease the value.

3 Rank three

At rank 3 the collapsing law $(x^y)^z = x^{yz}$ reduces (HT_r) to a comparison of exponents.

Lemma 3.1. *If $y \geq 2$ then $yz \leq y^z$ for every z . If moreover $z \geq 3$ the inequality is strict.*

Proposition 3.2. *If $y \geq 2$ then $yz = y^z$ if and only if $z = 1$ or $(y, z) = (2, 2)$.*

Proof. If $z = 0$ then $yz = 0 \neq 1 = y^z$. If $z = 1$ both sides are y . If $z = 2$ then $yz = y^z$ reads $2y = y^2$, i.e. $y = 2$ after cancelling $y > 0$. If $z \geq 3$ the inequality is strict by Lemma 3.1. The converse is immediate. $\square$

Proof of Theorem 1.1. Since $(x^y)^z = x^{yz}$, the inequality follows from Lemma 3.1 and monotonicity of $t \mapsto x^t$ for $x \geq 1$. For $x \geq 2$ the map $t \mapsto x^t$ is injective, so equality holds iff $yz = y^z$, and Proposition 3.2 applies. For the last claim, $(x^1)^z = x^z$ while $x^{(1^z)} = x$, and $x^z > x$ whenever $x, z \geq 2$. $\square$

Example 3.3 (Failure of strictness). Take $a = 2$ and $n = 4$. The five bracketings evaluate under H_3 to

$$((aa)a)a = 256, \quad (a(aa))a = 256, \quad (aa)(aa) = 256, \quad a((aa)a) = 65536, \quad a(a(aa)) = 65536.$$

In particular the Tamari cover $((aa)a)a \triangleleft (aa)(aa)$ has equal endpoints, so $\text{ev}_{3,2}$ is monotone but not strictly monotone; it is also not injective.

4 Rank four

At rank 4 we know of no identity that would play the role of $(x^y)^z = x^{yz}$, and a different argument is required. The proof below is by induction on z ; the base case $z = 1$ must be treated separately, since the inductive hypothesis at $z = 0$ is too weak to survive the step.

Lemma 4.1 (Squaring). *If $x \geq 2$ and $m \geq 1$ then $(x \uparrow\uparrow m)^2 \leq x \uparrow\uparrow (m + 1)$.*

Proof. Induction on m . For $m = 1$ the claim is $x^2 \leq x^x$, which holds as $x \geq 2$. For the step put $t = x \uparrow\uparrow m$; then $(x \uparrow\uparrow (m+1))^2 = (x^t)^2 = x^{2t} \leq x^{(x^t)} = x \uparrow\uparrow (m+2)$, where $2t \leq 2^t \leq x^t$ uses $t \geq 1$ and $x \geq 2$. $\square$

Remark 4.2. The direct inequality $t^2 \leq x^t$ is *false* in general — take $x = 2$, $t = 3$, where $t^2 = 9 > 8 = x^t$ — so the inductive formulation in Lemma 4.1 is essential.

Lemma 4.3. *If $y \geq 2$ and $m \geq 2$ then $m+2 \leq y^m$. If moreover $m \geq 3$ then $m+3 \leq y^m$.*

Proof of Theorem 1.2. Write $y = y' + 1$ with $y' \geq 1$ and induct on z . For $z = 0$ the left side is 1 and the right side is $x \geq 1$. For $z = 1$ both sides equal $x \uparrow\uparrow y$. Assume $z \geq 1$ and set $m := y \uparrow\uparrow z$, so $m \geq y \geq 2$ and $y' \leq m$. By Lemma 4.3 we have $y^m \geq m+2$, so $y^m = k+1$ with $k \geq m+1$. Writing $A := x \uparrow\uparrow y = x^{(x \uparrow\uparrow y')}$ we compute

$$\begin{aligned} (x \uparrow\uparrow y) \uparrow\uparrow (z+1) &= A^{(x \uparrow\uparrow y) \uparrow\uparrow z} \leq A^{x \uparrow\uparrow m} = (x^{x \uparrow\uparrow y'})^{x \uparrow\uparrow m} = x^{(x \uparrow\uparrow y')(x \uparrow\uparrow m)} \\ &\leq x^{(x \uparrow\uparrow m)^2} \leq x^{x \uparrow\uparrow (m+1)} \leq x^{x \uparrow\uparrow k} = x \uparrow\uparrow (k+1) = x \uparrow\uparrow (y \uparrow\uparrow (z+1)), \end{aligned}$$

using the inductive hypothesis, $x \uparrow\uparrow y' \leq x \uparrow\uparrow m$, Lemma 4.1, and $m+1 \leq k$. This proves the inequality.

For strictness, note that the penultimate step is strict as soon as $m \geq 3$: then $y^m \geq m+3$ by Lemma 4.3, hence $m+1 < k$ and $x \uparrow\uparrow (m+1) < x \uparrow\uparrow k$. Now $m = y \uparrow\uparrow z \geq 3$ holds whenever $z \geq 2$ (then $m \geq 4$), and also when $z = 1$ and $y \geq 3$ (then $m = y$). This gives strictness at $z+1$ for every $z+1 \geq 2$ except $(y, z+1) = (2, 2)$, which we compute directly:

$$(x \uparrow\uparrow 2) \uparrow\uparrow 2 = (x^x)^{x^x} = x^{x \cdot x^x} = x^{x^{x+1}}, \quad x \uparrow\uparrow (2 \uparrow\uparrow 2) = x \uparrow\uparrow 4 = x^{x^{x^x}},$$

and $x+1 < x^x$ for $x \geq 2$. Hence the inequality is strict for every $z \geq 2$. $\square$

5 The dichotomy

Proof of Corollary 1.5. Order-preservation at every rank $r \geq 3$ is Proposition 1.3. Non-strictness at rank 3 is Example 3.3. Strictness at rank 4 is Theorem 1.2; for general $r \geq 4$ it is Theorem 1.4, which is proved in Section 6 and rests on Theorem 1.7(3). (The present section proves only the rank-4 case, so that the dichotomy is available before the general induction.) $\square$

Proof of Proposition 1.3. Induction on the derivation of the cover. At the root the claim is Theorem 1.6 instantiated at the values of the three subtrees, which are all ≥ 2 because $a \geq 2$. If the rotation takes place inside the left subtree, one uses monotonicity of $H_r(\cdot, v)$ in its first argument; inside the right subtree, monotonicity of $H_r(u, \cdot)$ in its second argument, valid for $u \geq 2$. Both monotonicity statements hold for every $r \geq 1$ and are proved by a double induction on the rank and on the argument. The formula for the right comb is immediate from $H_{r+1}(a, n+1) = H_r(a, H_{r+1}(a, n))$ together with $H_{r+1}(a, 1) = a$. $\square$

The mechanism is worth stating plainly. At rank 3 the identity $(x^y)^z = x^{yz}$ turns (HT_r) into $yz \leq y^z$, an inequality which is *tight*: it becomes an equality at $(y, z) = (2, 2)$. Each of the three coincidences in Example 3.3 is an instance of this one equality, applied at a node whose two lower subtrees both evaluate to 2: the cover $((aa)a)a \leq (aa)(aa)$ is the root rotation with $(x, y, z) = (4, 2, 2)$, and the pairs $((aa)a)a$, $(a(aa))a$ and $a((aa)a)$, $a(a(aa))$ differ only in a subtree $(aa)a = a(aa) = 16$, which is the root rotation with $(x, y, z) = (2, 2, 2)$. At rank 4 Theorem 1.2 shows that no equality with $z \geq 2$ occurs at all, and the gap is already visible at the smallest instance: for $x = y = z = 2$ the two sides are 256 and 65536.

6 General rank

We now prove Theorem 1.6. The engine is the following inequality; write

$$H_{r-1}(H_r(x, a), H_r(x, b)) \leq H_r(x, a + b). \quad (\text{SUM}_r)$$

Proposition 6.1. *Let $r \geq 3$ and $x \geq 2$.*

1. *If (HT_{r-1}) holds, then (SUM_r) holds for all $a, b \geq 1$.*
2. *If (SUM_r) holds and $u + v \leq H_{r-1}(u, v)$ for all $u, v \geq 2$, then (HT_r) holds for all $y \geq 2$.*

Proof. (1) Induct on a . For $a = 1$ both sides are equal:

$$H_{r-1}(H_r(x, 1), H_r(x, b)) = H_{r-1}(x, H_r(x, b)) = H_r(x, b + 1).$$

For the step, $H_r(x, a + 1) = H_{r-1}(x, H_r(x, a))$, so by (HT_{r-1})

$$H_{r-1}(H_{r-1}(x, H_r(x, a)), H_r(x, b)) \leq H_{r-1}(x, H_{r-1}(H_r(x, a), H_r(x, b))),$$

and the inductive hypothesis bounds the inner term by $H_r(x, a + b)$, giving $H_{r-1}(x, H_r(x, a + b)) = H_r(x, a + b + 1)$.

(2) Induct on z , with $z = 0, 1$ as in Theorem 1.2. For the step set $m := H_r(y, z) \geq 2$; then $H_r(H_r(x, y), z + 1) = H_{r-1}(H_r(x, y), H_r(H_r(x, y), z)) \leq H_{r-1}(H_r(x, y), H_r(x, m)) \leq H_r(x, y + m) \leq H_r(x, H_{r-1}(y, m)) = H_r(x, H_r(y, z + 1))$. $\square$

Proof of Theorem 1.6. Induction on r . The base case $r = 3$ is Theorem 1.1. For the step, the inductive hypothesis is exactly (HT_{r-1}) , so Proposition 6.1(1) gives (SUM_r) , and Proposition 6.1(2) then gives (HT_r) , the hypothesis $u + v \leq H_{r-1}(u, v)$ for $u, v \geq 2$ being Lemma 6.2. $\square$

Lemma 6.2. *For $u, v \geq 2$ and $s \geq 2$ we have $u + v \leq H_s(u, v)$.*

Proof. $u + v \leq uv = H_2(u, v)$, and $H_2(u, v) \leq H_s(u, v)$ by monotonicity in the rank. $\square$

Strictness

Lemma 6.3. *For $u, v \geq 2$ with $u \geq 3$ or $v \geq 3$, and $s \geq 2$, we have $u + v < H_s(u, v)$. Equality $u + v = H_s(u, v)$ occurs exactly at $u = v = 2$, where both sides are 4.*

Proof. $u + v < uv$ holds iff $(u - 1)(v - 1) > 1$, i.e. iff $u \geq 3$ or $v \geq 3$; and $uv = H_2(u, v) \leq H_s(u, v)$. At $u = v = 2$ we have $u + v = 4 = H_2(2, 2) = H_s(2, 2)$ by Remark 1.8. $\square$

Proof of Theorem 1.7. (1) $H_r(H_r(x, y), 1) = H_r(x, y) = H_r(x, H_r(y, 1))$ since $H_r(a, 1) = a$. (2) is Theorem 1.1. For (3) we induct on $r \geq 4$, the base $r = 4$ being the strictness clause of Theorem 1.2. Let $r \geq 5$ and write $z = w + 1$ with $w \geq 1$, and $m := H_r(y, w) \geq 2$.

If $y \geq 3$ or $w \geq 2$, then $y \geq 3$ or $m \geq 4$, so Lemma 6.3 makes the last step of the proof of Proposition 6.1(2) strict, and strictness propagates because $H_r(x, \cdot)$ is strictly increasing.

The remaining case is $y = z = 2$. Put $A := H_r(x, 2) = H_{r-1}(x, x) \geq 4$. Then

$$H_r(H_r(x, 2), 2) = H_{r-1}(A, A), \quad H_r(x, H_r(2, 2)) = H_r(x, 4) = H_{r-1}(x, H_{r-1}(x, A)),$$

using $H_r(2, 2) = 4$ from Remark 1.8. Since $A = H_{r-1}(x, x)$, the required inequality

$$H_{r-1}(H_{r-1}(x, x), A) < H_{r-1}(x, H_{r-1}(x, A))$$

is the inductive hypothesis at rank $r - 1$ applied to (x, x, A) , and $A \geq 2$. $\square$

We can now upgrade Proposition 1.3 to its strict form.

Proof of Theorem 1.4. Induction on the derivation of the cover, exactly as for Proposition 1.3 but with every inequality replaced by its strict counterpart. At the root the claim is Theorem 1.7(3) instantiated at the values of the three subtrees, which are all ≥ 2 because $a \geq 2$ (so in particular the hypothesis $z \geq 2$ of that item is met). If the rotation takes place inside the left subtree, one needs

$$2 \leq u < u' \implies H_r(u, v) < H_r(u', v) \quad (v \geq 1),$$

the strict form of monotonicity in the first argument; if it takes place inside the right subtree, one needs

$$1 \leq v < v' \implies H_r(u, v) < H_r(u, v') \quad (u \geq 2, r \geq 2),$$

the strict form of monotonicity in the second argument. Both are proved by the same double induction on the rank and on the argument as their non-strict counterparts, and both apply here because every subtree value is ≥ 2 . $\square$

The auxiliary facts used above — monotonicity of $H_r(x, \cdot)$ in its second argument, monotonicity in the rank, and $n + 1 \leq H_r(x, n)$ for $x \geq 2$, $r \geq 2$, $n \geq 1$ — are elementary and are proved in the formalization by a sequence of inductions.

The hypothesis $y \geq 2$ of Theorem 1.6 is sharp, in the following precise sense: for every $r \geq 3$ and all $x, z \geq 2$, taking $y = 1$ reverses the inequality. Indeed $H_r(x, 1) = x$, so the left side is $H_r(x, z)$, whereas $H_r(1, z) = 1$ makes the right side $H_r(x, 1) = x$, and $H_r(x, z) > x$ because $z \geq 2$. The restriction to $x, z \geq 2$ is necessary and not an artefact: for $x = 1$ both sides equal 1, since $H_r(1, \cdot) = 1$; for $z = 1$ they are equal by Theorem 1.7(1); and for $z = 0$ the left side is 1 and the right side is x , so the original inequality still holds. For $x = 0$ the operations degenerate.

7 Arbitrary leaf values

Nothing so far has used the fact that the leaves all carry the same value. The local law (HT_r), and its strict form, are statements about three independent arguments; only the evaluation $\text{ev}_{r,a}$ was tied to a single label. Removing that restriction costs nothing and gives a strictly stronger statement, which is moreover the one that the discussion recalled in Section 9 was about.

Definition 7.1 (Labelled trees). A *labelled tree* is a binary tree each of whose leaves carries a natural number. Its evaluation at rank r is defined by $\text{ev}_r(\ell) = \ell$ for a leaf labelled ℓ and $\text{ev}_r(T \cdot U) = H_r(\text{ev}_r T, \text{ev}_r U)$. Right rotation at a node is defined exactly as before; it permutes no labels.

Theorem 7.2 (Arbitrary leaf values). *Let T be a labelled tree all of whose leaf labels are ≥ 2 , and let $T \triangleleft U$ be a Tamari cover. Then*

$$\text{ev}_r(T) \leq \text{ev}_r(U) \quad (r \geq 3), \quad \text{ev}_r(T) < \text{ev}_r(U) \quad (r \geq 4),$$

and the same holds along the Tamari order itself, that is along the reflexive transitive closure of the covering relation in the first case and its transitive closure in the second.

Proof. Rotation does not change the multiset of leaf labels, so the hypothesis is preserved along a cover, and every subtree of T evaluates to at least 2 at every rank $r \geq 2$ by induction on the tree. The induction on the derivation of the cover is then the one already used for

Proposition 1.3 and Theorem 1.4: at the root the claim is Theorem 1.6 respectively Theorem 1.7(3), instantiated at the three subtree values, which are independent of one another and all ≥ 2 ; inside the left subtree one uses monotonicity, respectively strict monotonicity, of $H_r(\cdot, v)$ in its first argument, and inside the right subtree the corresponding statement for $H_r(u, \cdot)$. The passage to the closures is by induction on the length of the chain. $\square$

We can now say exactly when a cover fails to be strict. At rank 3 the only non-trivial equality of (HT_r) is $(y, z) = (2, 2)$, and in a labelled tree whose labels are all ≥ 2 an *internal* node evaluates to at least $H_r(2, 2) = 4$; so a subtree evaluates to 2 exactly when it is a leaf labelled 2.

Corollary 7.3 (Equality along a cover). *Let all labels be ≥ 2 and consider the rotation $((A \cdot B) \cdot C) \rightarrow (A \cdot (B \cdot C))$ at rank 3. Its two sides are equal if and only if B and C are both leaves labelled 2. Consequently, for a single repeated label a , every cover is strict unless $(r, a) = (3, 2)$; and for $(r, a) = (3, 2)$ equality does occur, as Example 3.3 shows.*

Proof. By Theorem 1.1 the two sides agree exactly when $\text{ev}(C) = 1$ or $(\text{ev}(B), \text{ev}(C)) = (2, 2)$. The first is impossible because every label is ≥ 2 . For the second, a subtree evaluates to 2 only if it is a leaf labelled 2, by the remark above. For a single label a : if $a \geq 3$ no leaf is labelled 2, so no cover at rank 3 is an equality, and from rank 4 on Theorem 1.4 applies. $\square$

Thus the collapse at rank 3 is not merely rare but completely located: it happens at exactly one local configuration, two leaves labelled 2 being rotated past each other, and from rank 4 on it does not happen at all.

Setting every label equal to a recovers Proposition 1.3 and Theorem 1.4; the formalization carries out this specialisation explicitly — for covers, for the Tamari order, and in both the non-strict and the strict form — so the earlier statements are corollaries rather than separate results.

Theorem 7.2 is what extends the rank-3 discussion of Section 9 to the whole hierarchy. We emphasise again what it does not do: it does not decide when two incomparable bracketings can be separated in both directions, which is the half that was left open there.

8 The Lean development

The development is written for Lean 4 (v4.30.0) with Mathlib (v4.30.0) and is available at

<https://github.com/turahigashi/hyper-tamari>

It consists of the four files `Basic.lean`, `Tetration.lean`, `General.lean` and `VarLeaves.lean` in the directory `HyperTamari/`, together with the raw `lake build` output recording the axiom audit (`logs/axiom-audit.txt`).

The correspondence with the numbered statements of this paper is given in Table 1. The population of the table is fixed as follows: it consists of *every* numbered environment in Sections 1–7 — that is, every **theorem**, **proposition**, **lemma**, **corollary**, **definition**, **example** and **remark** environment in the source of those sections. There are 24 of them: 6 theorems, 3 propositions, 6 lemmas, 2 corollaries, 4 definitions, 1 example and 2 remarks. (Section 8 contains no numbered environments.) Every one of the 24 appears in the table, and the only other row is one recording the specialisation of the labelled-tree results to a single repeated label; both facts, together with the existence in the development of every declaration name printed in the table, are checked by a script run on the source of this paper and of the Lean files. A numbered statement with several clauses may occupy several rows. We distinguish three cases, because they carry different evidential weight:

EXACT a single declaration with the same quantifiers, hypotheses and conclusion as the printed clause;

COLLECTIVE the clause is covered by a combination of declarations (typically because the paper states several things at once);

PROSE background material that is *not* formalized here — standard facts about the Tamari lattice that we use as context, or explanatory remarks.

We stress the third row of this classification: the paper does *not* claim that everything it says is machine-checked. What is machine-checked is every clause of every numbered statement that asserts an inequality or an equality about H_r or about the evaluation of trees; the table says which declaration carries which clause.

Table 1: Correspondence between the 24 numbered statements of Sections 1–7 and the Lean development. Clauses of one statement are listed on separate rows.

Statement	Lean declaration(s)	Status
Thm. 1.1 (inequality)	ht_rank3	EXACT
Thm. 1.1 (equality iff)	ht_rank3_eq_iff	EXACT
Thm. 1.1 (sharpness of $y \geq 2$)	ht_reversed_at_y_one at $r = 3$ (universal), ht_needs_two_le_y (the witness $x = z = 3$)	COLLECTIVE
Thm. 1.2 (inequality)	tet_ht, ht_rank4	EXACT
Thm. 1.2 (strict for $z \geq 2$)	tet_ht_strict_of_two_le, ht_rank4_strict_two	EXACT
Prop. 1.3 (covers)	Rot.eval_le	EXACT
Prop. 1.3 (Tamari order)	Rot.eval_le_of_reflTransGen	EXACT
Prop. 1.3 (right comb)	rightComb, eval_rightComb (the Lean rightComb n has $n + 1$ leaves, so the closed form is stated with the index shifted by one)	COLLECTIVE
Thm. 1.4 (covers)	Rot.eval_lt	EXACT
Thm. 1.4 (Tamari order)	Rot.eval_lt_of_transGen	EXACT
Cor. 1.5	Rot.eval_le, rank3_rotation_not_strict, Rot.eval_lt	COLLECTIVE
Thm. 1.6	ht_general	EXACT
Thm. 1.7 (1)–(4)	equality_classification, whose four components are ht_eq_at_one, ht_rank3_equality_witness with ht_rank3_eq_iff for (2), ht_strict_general, ht_lt_at_zero	COLLECTIVE
Rem. 1.8 ($H_r(x, y) > x$; $H_r(2, 2) = 4$)	self_lt_hyper, hyper_two_two	COLLECTIVE
Rem. 1.8 (the two coordi- nates compete)	— (explanatory)	PROSE
Def. 7.1	LTree, LTree.eval, LTree.AllLeavesTwoLe, LRot	EXACT
Thm. 7.2	LRot.eval_le, LRot.eval_lt, LRot.eval_le_of_reflTransGen, LRot.eval_lt_of_transGen, LRot.preserves, LTree.two_le_eval	COLLECTIVE
Specialisation to a single la- bel	ofBTree, eval_ofBTree, rot_ofBTree, reflTransGen_ofBTree, transGen_ofBTree, Rot.eval_le_of_var, Rot.eval_lt_of_var, Rot.eval_le_of_reflTransGen_of_var, Rot.eval_lt_of_transGen_of_var	COLLECTIVE

Table 1 (continued)

Statement	Lean declaration(s)	Status
Cor. 7.3	LTree.four_le_eval_node, LTree.eval_eq_two_iff, LRot.root_eq_rank3_iff, LRot.root_lt_rank3_of_three_le, rank3_root_lt_of_three_le, cover_strict_unless_rank3_two	COLLECTIVE
Def. 2.1	hyper, hyper_zero, hyper_one_zero, hyper_two_zero, hyper_succ_succ_zero, hyper_succ	EXACT
Lem. 2.2	hyper_one, hyper_two, hyper_three, hyper_one_arg	COLLECTIVE
Def. 2.3 (trees, ev)	BTree, BTree.eval: an equivalent encoding, on all finite binary trees rather than on the fixed- arity $\mathcal{T}_n$	COLLECTIVE
Def. 2.4 (rotation, Tamari order)	Rot (cover), Relation.ReflTransGen Rot (reachability): an equivalent encoding; antisym- metry, and hence that this reachability relation is a partial order, is not formalized	COLLECTIVE
Def. 2.4 (lattice, comb ex- tremality)	— (classical [2, 12], used as background)	PROSE
Lem. 3.1	mul_le_pow_of_two_le, mul_lt_pow_of_three_le	COLLECTIVE
Prop. 3.2	mul_eq_pow_iff	EXACT
Ex. 3.3 (the five values; the colliding cover)	rank3_five_values, rank3_rotation_not_strict	COLLECTIVE
Lem. 4.1	sq_tet_le	EXACT
Rem. 4.2	sq_le_pow_fails (the negated universal state- ment), two_pow_three_lt_three_sq	EXACT
Lem. 4.3	add_two_le_pow, add_three_le_pow	COLLECTIVE
Prop. 6.1 (1), (2)	sum_lemma, ht_of_sum	COLLECTIVE
Lem. 6.2	add_le_hyper	EXACT
Lem. 6.3 (strict inequality; equality iff $u = v = 2$; value 4)	add_lt_hyper, add_eq_hyper_iff, hyper_two_two	COLLECTIVE

Two assertions of Section 6 are not numbered and therefore lie outside the population of the table; both are formalized: the sharpness statement for all ranks (“for every $r \geq 3$ and all $x, z \geq 2$, taking $y = 1$ reverses the inequality”) is `ht_reversed_at_y_one`, and the equality of both sides at $x = 1$ is `ht_eq_at_x_one`. The value pair $256 < 65536$ quoted in Section 5 is `ht_rank4_strict_smallest`. The development also contains `pow_lt_pow_right`’ and `pow_right_injective`’, choice-free replacements for the corresponding Mathlib lemmas (see below), and the auxiliary monotonicity facts `succ_le_hyper`, `hyper_mono_arg`, `hyper_mono_base`, `hyper_rank_mono`, `hyper_lt_arg` and `hyper_lt_base`.

Draft note — not part of the submission. The artifact has *not* yet been deposited in a long-term archive, so no DOI is quoted above. A version DOI will be minted and inserted here at the moment this paper is submitted, and the archive record will be linked to the submission. Until then the repository URL is the only pointer.

No declaration uses `sorry`, `native_decide` or any private axiom. Every one of the 116 theorems and lemmas in these four files is audited with `#print axioms`, and the audit is machine-checked to cover them all (declared 116, audited 116, missing 0; the script is `scripts/check_audit.py`); `sorryAx` does not occur. The axiom dependencies are `propext` alone (5 theorems), the pair `propext`, `Quot.sound` (105 theorems), and six theorems that

depend on no axioms at all. In particular

no declaration depends on `Classical.choice`.

This is not a foundational claim, and choice-freeness is not what the paper is about: every statement here is an elementary assertion about natural numbers, and nobody would expect it to need choice. We record it because it costs nothing and makes the artifact reusable — a development free of `Classical.choice` can be transported into constructive settings, or used inside a larger development that must itself avoid choice, without reproving anything.

An earlier version of the development did depend on `Classical.choice` in 30 declarations. The dependency was not in the mathematics but in four incidental sources — two Mathlib lemmas, the `nlinarith` tactic, and `norm_num` applied to a numeric inequality — and removing it changed no statement. The details are recorded in the repository.

9 Related work

Csákány and Waldhauser [1] introduce the associative spectrum and show that exponentiation is Catalan: for pairwise distinct prime leaves, distinct bracketings give distinct values. Their concern is the number of values; the order in which those values arrange themselves is not addressed, and the Tamari lattice appears in their bibliography only as the source of the enumeration of bracketings. Huang and Lehtonen [9] extend the invariant to the *associative-commutative spectrum* and determine it for exponentiation (their Proposition 6.1.5): the associative spectrum of exponentiation is C_{n-1} and its ac-spectrum is n^{n-1} . Their question is again how many distinct term operations the bracketings produce; ours is how the resulting values are *arranged*, and neither the Tamari order nor any monotonicity statement occurs in that line of work. Huang, Mickey and Xu [10] compute the associative spectrum of a further operation. Tamari [12] enumerated the bracketings, and Friedman and Tamari [2] proved that the right-rotation order on them is a finite lattice — their title already calls the relation a *loi demi-associative*, which is precisely the law we show to run through the whole hyperoperation hierarchy. Zeilberger [13] gives a sequent calculus for that law. Gryszka [8] studies the comparison of power towers by normalizing them, which concerns right-associated towers rather than general bracketings.

Where the rank-three case comes from, and what we add

It is important to say clearly what part of Proposition 1.3 is already known. In a MathOverflow question of January 2018 [3], Spivak treats exponentiation on the integers ≥ 2 as an ordered magma and orders bracketings by the *evaluation order*: $E \leq F$ when the inequality holds *for every* assignment of values to the leaves. He asks how this order relates to the height (Stanley) order on Dyck paths. In the answers, Chow proves that a Stanley cover does decrease the value, the computation resting on

$$B^{CD} = (B^C)^D \geq (BC)^D \geq BC^D \quad (B, C \geq 2, D \geq 1),$$

and Rubey reduces the same covering relation to the inequality $L_2^{(L_1^{xR})} \geq (L_2^{L_1})^{(xR)}$, an instance of rank-3 semi-associativity in a variable-leaf form.

Since the evaluation order quantifies over *all* leaf assignments it is stronger than the single-label evaluation $\text{ev}_{3,a}$, and since the Stanley lattice is an extension of the Tamari lattice, comparability in the Tamari order implies comparability in the Stanley order. Consequently **the rank-3 monotonicity is already contained, in a stronger form, in that discussion, and we do not claim it.**

That discussion is, however, explicitly unfinished, and it is confined to rank 3. Chow states that he can “complete half of the proof”, the half we have just quoted; the converse half — that incomparable bracketings can be separated in both directions by a suitable choice of leaf values — he leaves open, writing that he has “not quite figured out how to say this rigorously”. Rubey likewise reduces the covering relation to the displayed inequality and stops there. Neither treats any rank above 3, and neither asks when equality holds.

What Theorem 7.2 above adds is the extension from exponentiation to the whole Goodstein hierarchy of the *restriction to Tamari rotations* of that monotonicity direction: for arbitrary leaf values ≥ 2 the evaluation is monotone along the Tamari order at every rank $r \geq 3$, and *strictly* monotone from rank 4 on. We extend the rank; we do not extend the order, since the statement proved in that discussion is about the larger Stanley order. We do not settle Chow’s converse half, which remains open.

There is also a body of peer-reviewed work on orders attached to iterated exponentials, and it is worth saying precisely how it relates to the present question, because the titles suggest a closer overlap than there is. Brunson [4] fixes $e \leq x_1 \leq \dots \leq x_n$ and, *for each permutation* of the x_i , forms the tower $T(y, z, \dots, w) = y^{z^{\dots^w}}$, in which “association is always to the upper right”; the resulting partial order lives on the symmetric group S_n . Stembridge [11] and Griggs–Wachs [6] compare that order with the Bruhat order on S_n , the latter exhibiting a counterexample at $n = 5$ to their coincidence. Griggs [5] studies the poset of towers over the two-letter alphabet $\{a, b\}^n$, the tower again being “evaluated in the conventional way, from top down”.

In every one of these the *bracketing is fixed* — it is the right comb — and what varies is the assignment of entries to its leaves. In the single-label problem of Sections 2–6 the situation is exactly the reverse: the entries are fixed, all equal to a , and what varies is the bracketing. (Section 7 then lets the entries vary as well, but still over *all* bracketings rather than over the right comb alone.) The two families of questions are therefore independent: neither the Tamari order on bracketings nor any inequality of the form (HT_r) occurs in that line of work, whose partial orders live on S_n or on $\{a, b\}^n$ rather than on trees.

What we do claim

Not finding a result is not a proof that it is new, so we state the claim in pieces. The rank-3 monotonicity is known (above). For the following we have not found a prior occurrence: the local semi-associative law (HT_r) for *every* rank $r \geq 3$ of the Goodstein hierarchy; the complete classification of its equality cases, with the isolated rank-3 exception $(y, z) = (2, 2)$; the resulting dichotomy, strict on every Tamari cover from rank 4 on; and the machine-checked formalization.

Acknowledgements

Use of AI assistance

Artificial-intelligence assistants were used substantially in the preparation of this work, and the author sets that use out in detail here. No AI system is an author. The overall direction of the work and the choice of what to claim are the author’s, and the author is responsible for all of the submitted content.

Systems used. Anthropic’s Claude (Opus 5, through the Claude Code command-line interface) and OpenAI’s ChatGPT (Pro), during August and September 2026.

Use as part of the research method. The question studied was formulated by the author. Language models were then used to search for the shape of the proofs, to write the whole of the Lean 4 development, and to search the literature. Further language models were used as adversarial readers of successive drafts and of the formalization; their criticism led to the present form of the main theorem, in particular to the classification of the equality cases in Theorem 1.7, to the extension to arbitrary leaf values in Section 7, and to the complete localisation of the rank-3 collapse in Corollary 7.3. Two defects in the verification scripts themselves — a population count in Table 1 and a parser that silently skipped two declarations — were found this way and are fixed in the released version.

Use in preparing the manuscript. Language models drafted this text, including the abstract and the expository passages, from the author’s outline and from the formal statements. Every sentence was read and, where necessary, rewritten by the author before submission. All model output was treated as an untrusted suggestion.

What the author verified personally. The choice of definitions and conventions (Definition 2.1 in particular); the correspondence, statement by statement, between each informal claim of Sections 1–7 and the Lean declaration named for it in Table 1; the bibliographic data of every reference, against the publishers’ records; and the reproduction of the build and of the axiom audit from the released sources, with Lean 4 v4.30.0 and Mathlib v4.30.0 at revision c5ea00351c28 as pinned in the repository. No confidential third-party material — no unpublished manuscript and no referee report — was supplied to any external AI system.

Record. The Lean development, the source of this paper, the raw build log containing every `#print axioms` result, and the scripts that check the audit and the correspondence table are all in the accompanying repository, whose commit history records the development.

Draft note — not part of the submission. The journal’s policy requires that, where AI has assisted with a proof or an argument, the author verify every step personally and also verify every citation and attribution. When that has been done, the following sentence belongs here, and it must not be included before then: *The author has independently verified every step of the mathematical arguments printed in this paper, and every citation and attribution.* The bibliographic data have already been checked against the publishers’ records; what remains is to read the three items still marked [†] in the original, and to check the printed arguments by hand.

Two independent checks stand behind the mathematical content. Every clause of every numbered statement of Sections 1–7 that asserts an inequality or an equality about H_r or about the evaluation of trees has been formalised and checked by the Lean 4 kernel; Table 1 says, clause by clause, which declaration carries it, and marks the places where the paper relies instead on classical background about the Tamari lattice or on explanatory prose. The informal proofs printed in the text are the author’s rendering of the formal ones and are not themselves machine-checked; they are set out in full so that a reader can follow them without consulting the Lean sources.

It should be said precisely what the second of these checks does and does not establish. Conditional on the correctness of the formal statements, on the pinned versions of Lean and Mathlib, and on the released sources’ checking successfully, the validity of the proofs of the formalized statements is reduced to the Lean 4 kernel and the trusted parts of the toolchain. It is not thereby certified that a formal statement says what the corresponding English sentence says; that the definitions chosen are the natural or the standard ones; that Table 1 is complete; that the novelty claims and attributions are correct; or that the checking scripts are themselves right. Each of these rests on judgement and on the author’s verification, and is stated here so that it can be examined.

Draft note — not part of the submission. Items marked [†] in the list below have *not* been consulted in the original at the time of writing; the bibliographic data were taken from secondary sources and independently cross-checked against Crossref, but the contents were not read. They are: [FT] Friedman–Tamari 1967, cited only for the origin of the lattice and of the term *loi demi-associative* (the phrase occurs in the title itself); [Gry] — *now consulted in the original* (Lith. Math. J. **62**(2), 192–206, 15 pages): it compares right-associated power towers with varying entries and heights and does not mention bracketings, binary trees, Catalan structures, rotations or any lattice, which confirms the description given in Section 9; [Tamari] Tamari 1962, cited only for the enumeration of bracketings and for the two combs being the extreme elements; and [Ste] Stembridge 1989, on Brunson’s order, cited only to place the question among its peer-reviewed neighbours. [Bru], [Gri] and [GW] were also identified late, in response to a referee report, but have since been consulted in the original and are no longer marked: all three fix the bracketing to the right comb and vary the entries, as now described in Section 9. None of the results of the present paper depends on these three items: they are cited for attribution and context, and every mathematical statement used in the proofs is either proved here or formalised in the accompanying Lean development. Before submission the three still-unchecked items [FT], [Tamari] and [Ste] must be checked in the original, and this note removed by setting `\draftfalse`.

References

- [1] B. Csákány and T. Waldhauser, *Associative spectra of binary operations*, Multiple-Valued Logic **5** (2000), 175–200; arXiv:1102.2124.
- [2] [†] H. Friedman and D. Tamari, *Problèmes d’associativité: une structure de treillis finis induite par une loi demi-associative*, Journal of Combinatorial Theory **2** (1967), no. 3, 215–242.
- [3] D. Spivak, *Is the order on repeated exponentiation the Dyck order?*, MathOverflow question 289708 (1 January 2018), with answers by T. Chow and M. Rubey, <https://mathoverflow.net/q/289708>.
- [4] B. W. Brunson, *The partial order of iterated exponentials*, American Mathematical Monthly **93** (1986), 779–786.
- [5] J. R. Griggs, *Iterated exponentials of two numbers*, Discrete Mathematics **88** (1991), 193–209.
- [6] J. R. Griggs and M. L. Wachs, *Towers of powers and Bruhat order*, European Journal of Combinatorics **13** (1992), 367–370.
- [7] R. L. Goodstein, *Transfinite ordinals in recursive number theory*, Journal of Symbolic Logic **12** (1947), 123–129.
- [8] K. Gryszka, *How to compare power towers?*, Lithuanian Mathematical Journal **62** (2022), no. 2, 192–206.
- [9] J. Huang and E. Lehtonen, *The associative-commutative spectrum of a binary operation*, Discrete Mathematics **346** (2023), 113468; arXiv:2202.11826.
- [10] J. Huang, M. Mickey and J. Xu, *The nonassociativity of the double minus operation*, arXiv:1710.05651 (2017).
- [11] [†] J. R. Stembridge, *The Bruhat order and iterated exponentials*, Journal of Combinatorial Theory Series A **50** (1989), 87–99.

- [12] [†] D. Tamari, *The algebra of bracketings and their enumeration*, Nieuw Archief voor Wiskunde (3) **10** (1962), 131–146.
- [13] N. Zeilberger, *A sequent calculus for a semi-associative law*, Logical Methods in Computer Science **15** (2019), no. 1, 9:1–9:23.